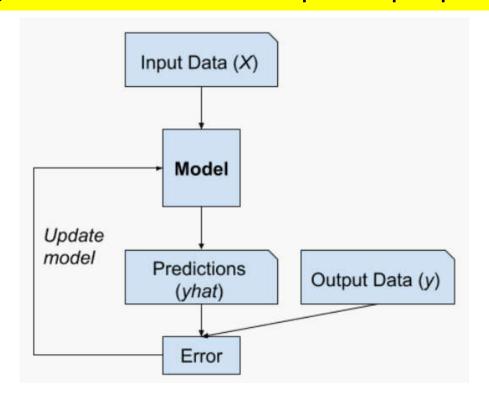
GENERATIVE ADVERSARIAL NETWORKS

References:

- Generative Adversarial Networks, Goodfellow et al.
- Conditional Generative Adversarial Nets
 Mehdi Mirza et al.
- Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, Zhu et al.
- Generative Adversarial Networks lecture (by Hongyang Zhang)

Supervised vs. Unsupervised Learning

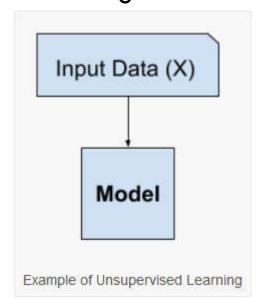
In the predictive methods, the goal is to learn a mapping from inputs x to outputs y, given a labeled set of input-output pairs



Supervised vs. Unsupervised Learning

unsupervised learning: only given inputs, and the goal is to find "interesting patterns" in the data.

clustering and generative modeling



- Generative modeling is an unsupervised learning task in machine learning
- automatically learning the patterns in input data
- * model used to generate new examples that plausibly drawn from the original dataset.
- In generative modeling, we'd like to train a distribution:
- such a way that the model generate new examples that plausibly could have been drawn from the original dataset.
- good generative model :generate new examples not just plausible, but indistinguishable from real examples

- ► Naïve Bayes
- **►**GMM
- ▶ Deep learning methods can be used as generative models
- ► A modern example of deep learning generative modeling algorithms
- we'd like to train a network that models a distribution
- ► For example, LLMs want to model the distribution of natural language
- ▶ We want to design a generative model to generate images.



- Given training data $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\} \sim p_{data}(\mathbf{x})$, the true data density
- Parameterize $p_{\theta}(\mathbf{x})$, the data density estimated by model
- Estimate θ by minimizing some "distance" between p_{data} (the unknown data density) and p_{θ} :

$$\min_{\boldsymbol{\theta}} \ \mathsf{dist}(p_{data} \| p_{\boldsymbol{\theta}})$$

- After training, can generate new data $\mathbf{x} \sim p_{\boldsymbol{\theta}}(\mathbf{x})$
- Need a training set from p_{data} and an explicit form of p_{θ}

Push-forward

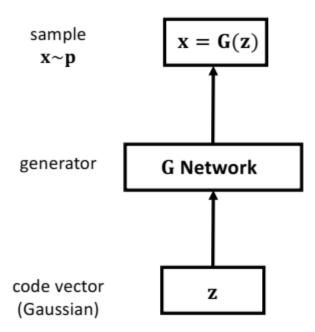
Theorem: Representation through push-forward

Let r be any continuous distribution on \mathbb{R}^h . For any distribution p on \mathbb{R}^d , there exist push-forward maps $G: \mathbb{R}^h \to \mathbb{R}^d$ such that

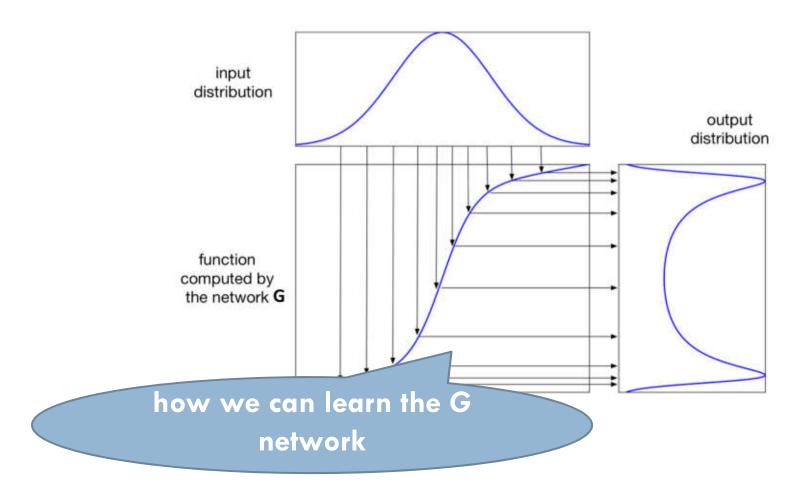
$$\mathbf{z} \sim r \implies \mathbf{G}(\mathbf{z}) \sim p$$

simply take r to be standard Gaussian noise

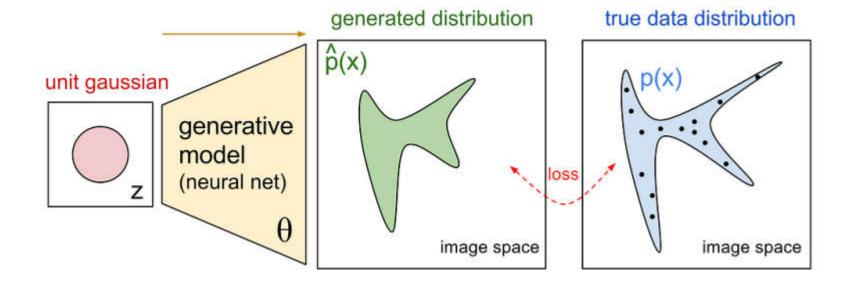
- Start by sampling the code vector z from a simple distribution (e.g., Gaussian)
- GAN computes a differentiable function G mapping z to an x in data space
- G maps one distribution to another



example

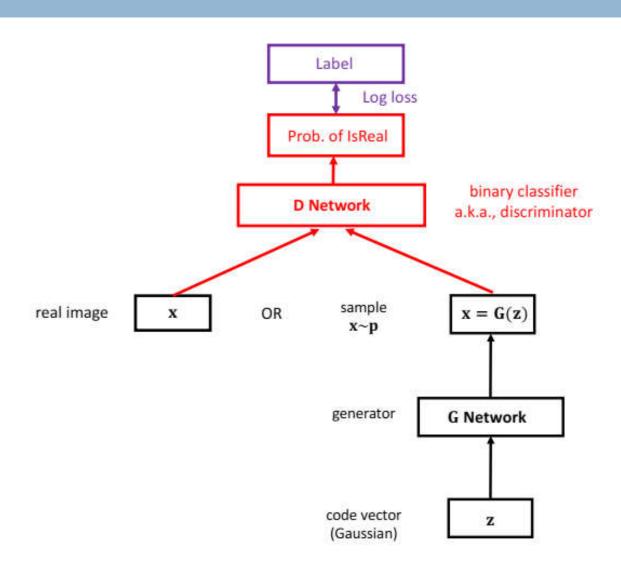


Learn Generator

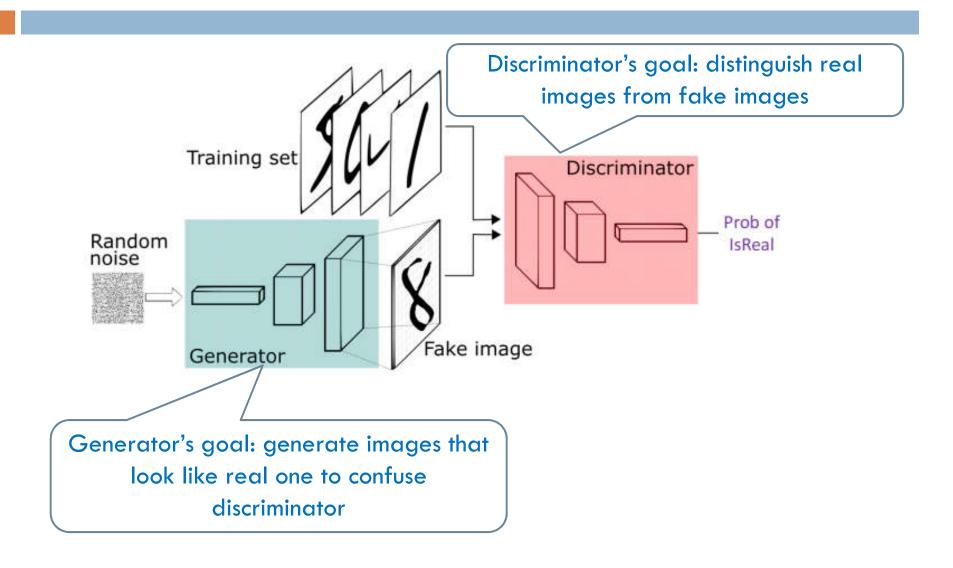


- How can we define the loss to distinguish two distributions?
 - Use a discriminator

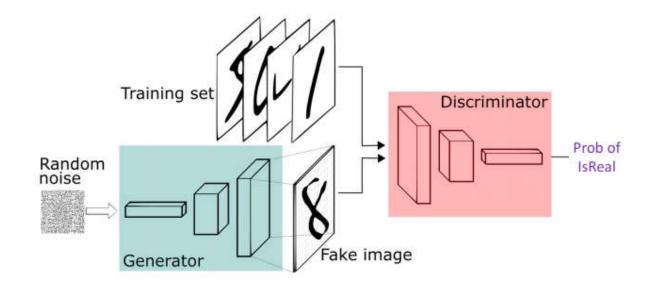
Generative Adversarial Networks structure



Generative Adversarial Networks



Discriminator's Loss

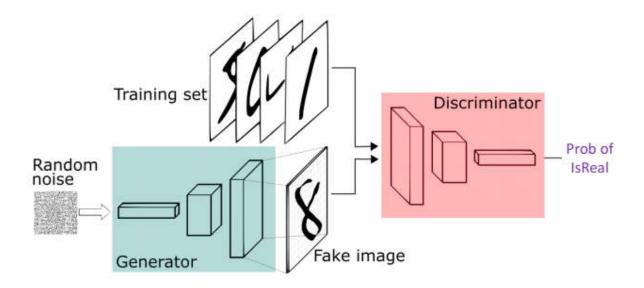


For a fixed generator G

- ▶ If x is real, minimize $-\log D(x)$; if x is fake, minimize $-\log(1 D(x))$
- Assume that x being from real/fake distribution is with equal chance:

$$\underbrace{\min_{D} \underbrace{-\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{data}}[\log D(\mathbf{x})]}_{\mathbf{x} \text{ is real with one half prob}} \underbrace{-\frac{1}{2} \mathbb{E}_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})}[\log (1 - D(G(\mathbf{z})))]}_{\mathbf{x} \text{ is fake with another half prob}}$$

Generator's Loss

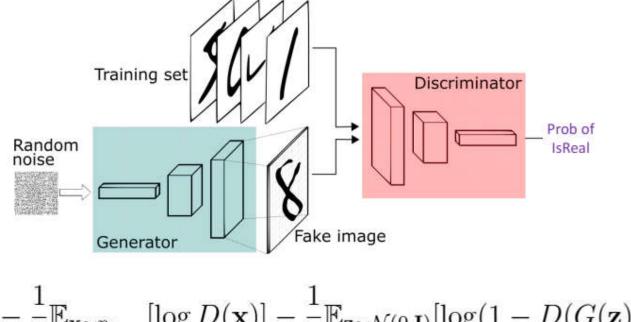


For a fixed Discriminator D

if x is fake, maximize $-\log(1 - D(\mathbf{x}))$

$$\max_{G} - \frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{data}}[\log D(\mathbf{x})] - \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})}[\log(1 - D(G(\mathbf{z})))]$$

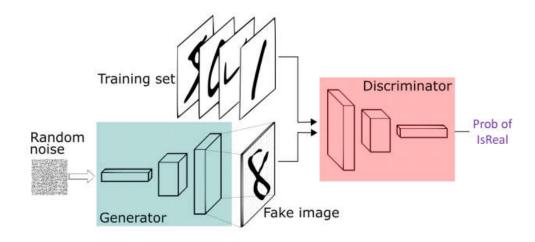
GAN



$$\max_{G} \min_{D} - \frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{data}} [\log D(\mathbf{x})] - \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})} [\log (1 - D(G(\mathbf{z})))]$$

$$\min_{G} \max_{D} \hat{\mathbb{E}}_{\mathbf{x} \sim p_{data}}[\log D(\mathbf{x})] + \hat{\mathbb{E}}_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})}[\log(1 - D(G(\mathbf{z})))]$$

How to solve



$$\min_{G} \max_{D} V(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{data}}[\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})}[\log(1 - D(G(\mathbf{z})))]$$

- Solved by alternative minimization-maximization:
 - ightharpoonup G step: Fix D and update G by one-step gradient descent
 - ightharpoonup D step: Fix G and update D by one-step gradient ascent
 - Repeat until the algorithm reaches an approximate equilibrium

Let $p_g(\mathbf{x})$ be the density of \mathbf{x} estimated by the generator G. For G fixed, the optimal discriminator D is $D_G^*(\mathbf{x}) = \frac{p_{data}(\mathbf{x})}{p_{data}(\mathbf{x}) + p_g(\mathbf{x})}$.

$$V(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{data}}[\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})}[\log(1 - D(G(\mathbf{z})))]$$

$$= \int_{\mathbf{x}} p_{data}(\mathbf{x}) \log D(\mathbf{x}) d\mathbf{x} + \int_{\mathbf{z}} p_{\mathbf{z}}(\mathbf{z}) \log(1 - D(G(\mathbf{z}))) d\mathbf{z}$$

$$= \int_{\mathbf{x}} \underbrace{p_{data}(\mathbf{x}) \log(D(\mathbf{x})) + p_{g}(\mathbf{x}) \log(1 - D(\mathbf{x}))}_{f(D(\mathbf{x}))} d\mathbf{x}.$$

$$D_{G}^{*}(\mathbf{x}) := \operatorname{argmax}_{D(\mathbf{x})} f(D(\mathbf{x})) = \underbrace{\frac{p_{data}(\mathbf{x})}{p_{data}(\mathbf{x}) + p_{g}(\mathbf{x})}}_{p_{data}(\mathbf{x}) + p_{g}(\mathbf{x})}.$$

Theorem 1: solution of G^*

The global minimum of $\min_{G} \max_{D} V(G, D)$ is achieved if and only if $p_g = p_{data}$. The optimal objective value is $-\log 4$.



Therefore, the generator is able to learn the data distribution p_{data} exactly if we can solve $\min_{G} \max_{D} V(G, D)$ exactly.

$$V(G, D_G^*) = \mathbb{E}_{\mathbf{x} \sim p_{data}} [\log D_G^*(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})} [\log (1 - D_G^*(G(\mathbf{z})))]$$

$$= \mathbb{E}_{\mathbf{x} \sim p_{data}} [\log D_G^*(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim p_g} [\log (1 - D_G^*(\mathbf{x}))]$$

$$= \mathbb{E}_{\mathbf{x} \sim p_{data}} \left[\log \frac{p_{data}(\mathbf{x})}{p_{data}(\mathbf{x}) + p_g(\mathbf{x})} \right] + \mathbb{E}_{\mathbf{x} \sim p_g} \left[\log \frac{p_g(\mathbf{x})}{p_{data}(\mathbf{x}) + p_g(\mathbf{x})} \right].$$

Recall that
$$KL(P||Q) = \mathbb{E}_{\mathbf{x} \sim P} \left[\log \frac{p(\mathbf{x})}{q(\mathbf{x})} \right]$$

$$V(G, D_G^*) = -\log 4 + KL\left(p_{data} \left| \left| \frac{p_{data} + p_g}{2} \right| \right) + KL\left(p_g \left| \left| \frac{p_{data} + p_g}{2} \right| \right) \right|$$

Jensen-Shannon divergence

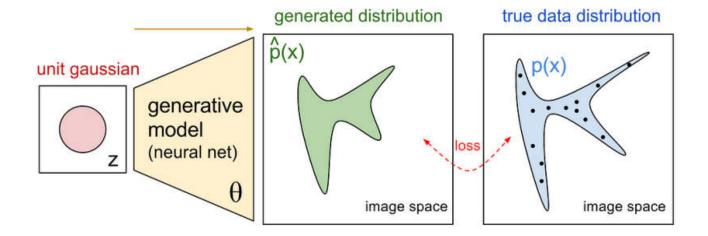
$$\mathrm{JSD}(P \parallel Q) = rac{1}{2}D(P \parallel M) + rac{1}{2}D(Q \parallel M),$$

where $M=rac{1}{2}(P+Q)$ is a mixture distribution of P and Q.

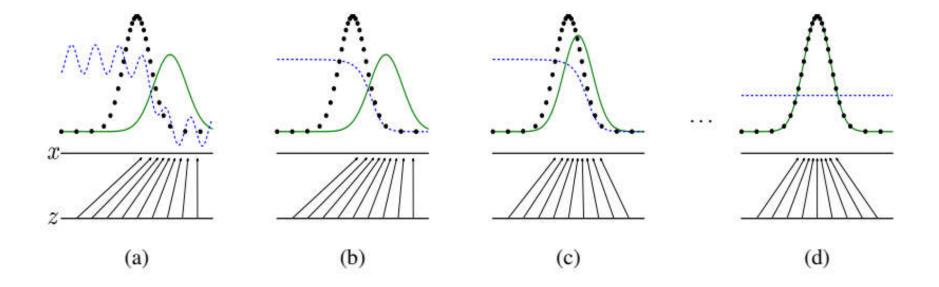
$$V(G, D_G^*) = -\log 4 + KL\left(p_{data} \left| \left| \frac{p_{data} + p_g}{2} \right| + KL\left(p_g \left| \left| \frac{p_{data} + p_g}{2} \right| \right. \right) \right.$$

$$= -\log 4 + 2 \cdot JSD(p_{data}||p_g) \ge -\log 4,$$

$$p_{data} = p_g$$
.



 Thus GAN is minimizing the Jensen–Shannon divergence between generated and real data distributions.



Conditional GAN

- Generative adversarial nets can be extended to a conditional model
- generator and discriminator are conditioned on some extra information y

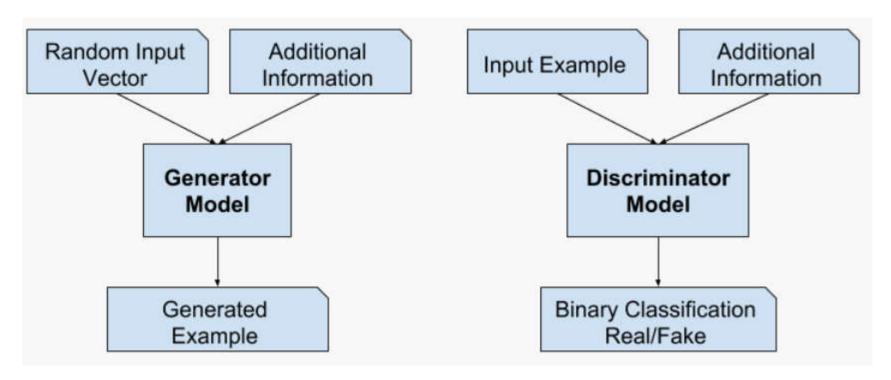


Image to Image Translation

- generating a new version of a given image with a specific modification, such as translating a summer landscape to winter.
- Training a model, a large dataset of paired examples

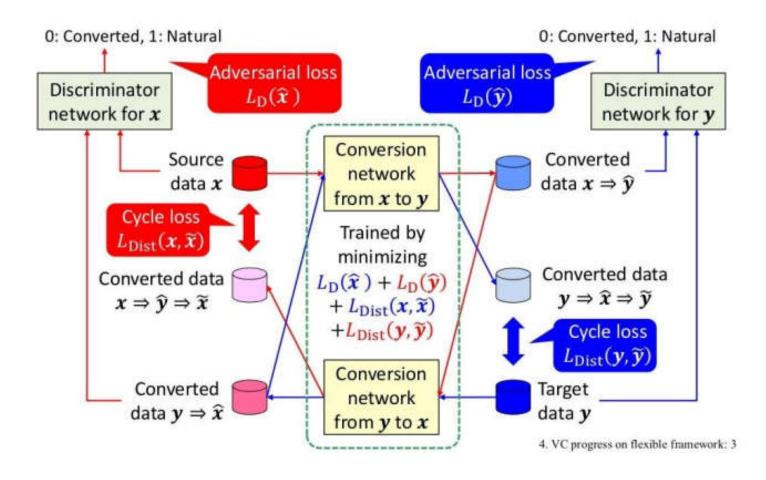


CycleGAN: Image to Image Translation

- CycleGAN: automatic training of image-to-image translation models without paired examples
- Unsupervised
- * A collection of images from the source and target domain
- not need to be related
- Extension of GAN
- Two generators, Two discriminators
- translation should be "cycle consistent"

Cycle Consistency

if we translate a sentence from English to Persian, and then translate it back from Persian to English, we should arrive back at the original sentence



- *translating images from summer to winter and winter to summer.
- *two databases of images, unpaired
- different locations at different times;

❖ Database 1: summer

❖ Database 2: winter

- ♦ two GANs
- *each GAN: a discriminator and a generator

•Generator Model 1:

Input: Images of summer Output Images of winter

•Discriminator Model 1:

- Input: Images of winter, output from Generator Model 1.
- Output: Probability of image is from database2

•Generator Model 2:

- Input: Images of winter Output: Images of summer
- Discriminator Model 2:
 - Input: Images of summer(database 1), output from Generator Model 2.
 - Output: Probability of image is from database 1.

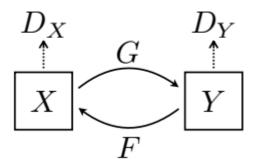
•Forward Cycle Consistency Loss:

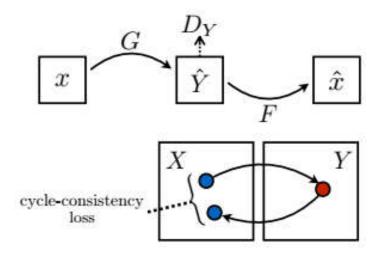
- Input summer images (DB1) to GAN 1
- Output winter images from GAN 1
- Input winter images from GAN 1 to GAN 2
- Output summer images from GAN 2
- Compare summer images (DB1) to summer images from GAN 2

•Backward Cycle Consistency Loss:

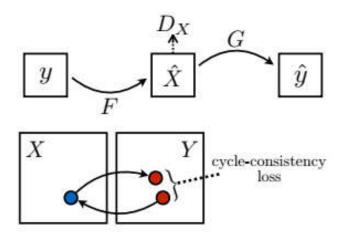
- Input winter images (DB 2) to GAN 2
- Output summer images from GAN 2
- Input summer images from GAN 2 to GAN 1
- Output winter images from GAN 1
- Compare winter images (DB 2) to winter images from GAN 1

 $G: X \to Y \text{ and } F: Y \to X$





$$x \to G(x) \to F(G(x)) \approx x$$



$$y \to F(y) \to G(F(y)) \approx y$$

CycleGAN Model Loss

$$G:X\to Y$$

$$\mathcal{L}_{GAN}(G, D_Y, X, Y) = \mathbb{E}_{y \sim p_{data}(y)}[\log D_Y(y)]$$

$$+ \mathbb{E}_{x \sim p_{data}(x)}[\log(1 - D_Y(G(x)))]$$

 $\min_{G} \max_{D_Y} \mathcal{L}_{GAN}(G, D_Y, X, Y)$

$$F: Y \to X$$

 $\min_F \max_{D_X} \mathcal{L}_{GAN}(F, D_X, Y, X)$

CycleGAN Model Loss

$$\mathcal{L}_{\text{cyc}}(G, F) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1].$$

$$\mathcal{L}(G, F, D_X, D_Y) = \mathcal{L}_{GAN}(G, D_Y, X, Y)$$

$$+ \mathcal{L}_{GAN}(F, D_X, Y, X)$$

$$+ \lambda \mathcal{L}_{cyc}(G, F),$$

$$G^*, F^* = \arg\min_{G, F} \max_{D_x, D_Y} \mathcal{L}(G, F, D_X, D_Y)$$